

HARSH VASISHT

+91 9650143510 ♦ harryvasisht@gmail.com ♦ [Linkedin](#) ♦ www.harsh.ai.in ♦ [Certificates](#) ♦ [GitHub](#)

TECHNICAL SKILLS

AI & ML	RAG, LLMs, Agentic AI, TensorFlow, PyTorch, Deep Learning, Computer Vision, NLP, Model Deployment, Fine-tuning, Prompt Optimization
Cloud & DevOps	GCP (Foundation Models), AWS, Docker, Kubernetes, Vector Databases, LangChain, Orchestration Frameworks
Programming	Python, JavaScript, R, SQL
Databases	MongoDB, Pinecone (Vector DB), PostgreSQL
Frameworks	Django, Flask, FastAPI, Power BI, Tableau, Grafana

EXPERIENCE

Associate AI Engineer Techolution	DEC 2023 – PRESENT <i>Hyderabad</i>
--------------------------------------	--

- Spearheaded enterprise-level Retrieval-Augmented Generation (RAG) system architecture for flagship internal and client-facing LLM applications, successfully processing 10M+ lines of code across 15+ programming languages and establishing the foundation for the company's GenAI product suite
- Pioneered and led cross-functional GenAI initiatives across multiple business units, orchestrating GCP Cloud infrastructure, Vector Databases, LangChain, and custom orchestration frameworks that became the standard for scalable AI solutions company-wide
- Revolutionized legacy code modernization through AI automation, leveraging GCP's Foundation Models (CodeLlama, Codey) to process 1000+ critical files with 95% accuracy, directly contributing to 500K+ in operational cost savings
- Engineered breakthrough AI-powered development workflows that transformed the entire development lifecycle, reducing manual effort by 70% across 25+ client projects and setting new industry benchmarks for AI-assisted software development
 - Achieved unprecedented model performance optimization, enhancing accuracy by 25% and inference speed by 40% through innovative prompt engineering and fine-tuning methodologies, making real-time AI applications feasible for enterprise clients
- Architected and deployed AI-powered code quality pipelines integrated into workflows, improving code maintainability by 80% and establishing new quality standards that reduced production bugs by 65%

GenAI Intern Attack Capital	OCT 2023 – NOV 2023 <i>USA (Remote)</i>
--------------------------------	--

- Architected cutting-edge conversational AI system by seamlessly integrating Open Source LLM agents (Llama, Mistral) and Closed Source models (GPT-3.5, GPT-4) with Twilio infrastructure, enabling automated call processing for 500+ daily high-value client interactions
- Engineered sophisticated Pinecone-based Knowledge Base infrastructure for intelligent call data management and real-time visualization, achieving 90% faster information retrieval and transforming the efficiency of AI calling agents
- Developed intelligent call routing and response system with advanced context-awareness capabilities, achieving 85% customer satisfaction and significantly improving client retention rates through personalized conversational experiences

EDUCATION

Bachelor of Technology in Artificial Intelligence, Gautam Buddha University, CGPA: 8.98 2020 - 2024

CERTIFICATIONS

Oracle Cloud Infrastructure Data Science Professional	2023
- Oracle Cloud Certificate Link	
Machine Learning with H2O Flow	2023
- Coursera Certificate Link	
IBM ML0101EN: Machine Learning with Python	2023
- Cognitive Class Certificate Link	

SELECTED PROJECTS

-
- THRIVE - Precision Care Challenge 2024 Hackathon**
AI & Backend Developer
June 2024
 - Engineered an **AI-powered patient triage system**, reducing unnecessary doctor visits by 30% and increasing triage accuracy by 65%.
 - Developed and deployed a state-of-the-art AI engine leveraging **MedQA, PubMedQA, MedMCQA, and MMLU** in a vector database.
 - Tech Stack:** Next.js, Supabase, Flask, MongoDB, AI-driven self-learning models
 - Deployed URL:** <https://med-ai-prod.vercel.app/>
 - Smart Coffee Machine (STMicroelectronics Partnership)**
Completed: 2023
 - Built contactless vital monitoring system** using PPG signals to track heart rate and caffeine intake correlation, achieving 94% accuracy with Arduino integration.
 - Deployed for 200+ users**, providing real-time health insights to track the impact of caffeine consumption.
 - Tech Stack:** Python, TensorFlow, Arduino, IoT, Real-time Processing
 - Project Link:** [GitHub Repository](#)